

Github1

February 9, 2024

1 1. Data Wrangling

```
[3]: import pandas as pd
import numpy as np
```

```
[4]: df=pd.read_csv('health care diabetes.csv')
```

```
[5]: df.shape
```

```
[5]: (768, 9)
```

```
[6]: df.head()
```

```
[6]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[7]: df.columns
```

```
[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
        'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
        dtype='object')
```

```
[8]: df.isna().sum()
```

```
[8]: Pregnancies      0
Glucose             0
```

```
BloodPressure      0
SkinThickness      0
Insulin            0
BMI                0
DiabetesPedigreeFunction  0
Age                0
Outcome            0
dtype: int64
```

```
[9]: #inference : the dataset contains no null values
```

```
[10]: #check for duplicated values
df.duplicated().sum()
```

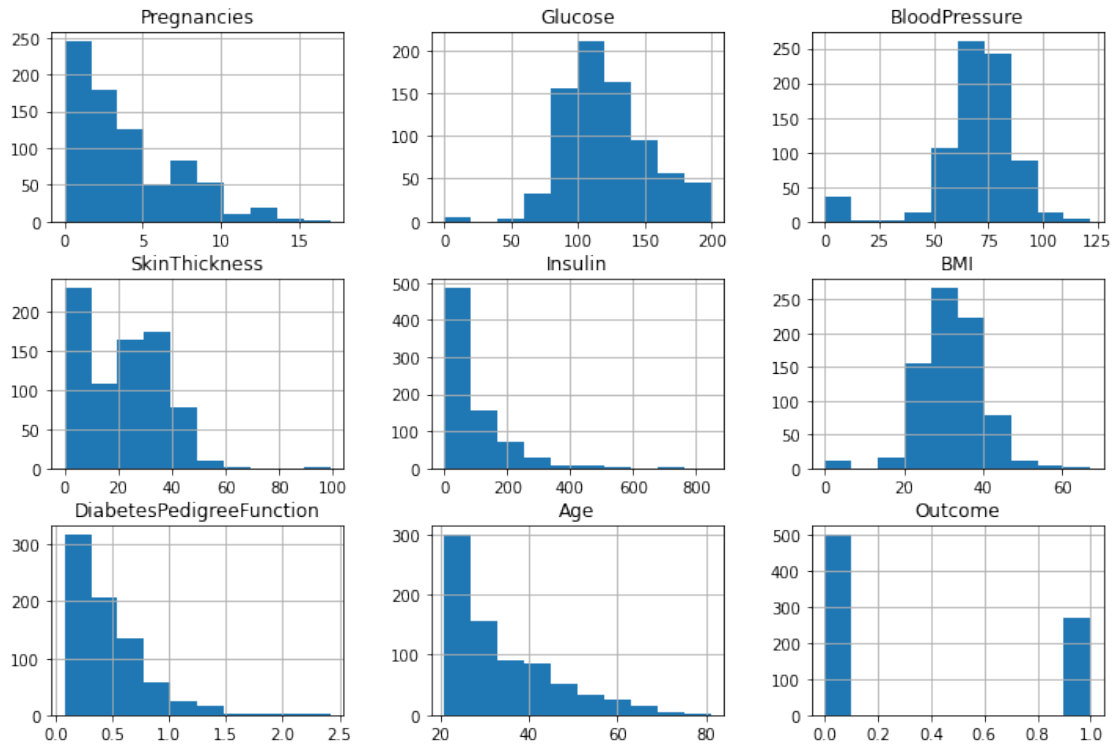
```
[10]: 0
```

```
[ ]: #inference : the dataset contains no duplicate entries
```

```
[11]: import matplotlib.pyplot as plt
```

```
[12]: df.hist(bins=10, figsize=(12, 8))
plt.suptitle('Histograms for All Columns', x=0.5, y=1.02, fontsize=16)
plt.show()
```

Histograms for All Columns



```
[13]: zeros_count_per_column = df.eq(0).sum()

print("Number of zeros in each column:")
print(zeros_count_per_column)
```

```
Number of zeros in each column:
Pregnancies          111
Glucose              5
BloodPressure        35
SkinThickness       227
Insulin             374
BMI                 11
DiabetesPedigreeFunction  0
Age                 0
Outcome            500
dtype: int64
```

The zero entries represent null value. These zero entries are to be replaced with mean or mode.
 ##### Assumption: As the gender of patients is not given, it is assumed that zero values in Pregnancies represents null value of female patients only.

2 Ques 2(b)

```
[14]: # Calculate the column-wise mean excluding zeros
column_means = df[df != 0].replace(0, np.nan).mean()

# Replace zeros with the column mean (excluding 'Outcome')
df[df.columns.difference(['Outcome'])] = df[df.columns.difference(['Outcome'])].
    ↪replace(0, np.nan)
df = df.fillna(column_means)

print(df)
```

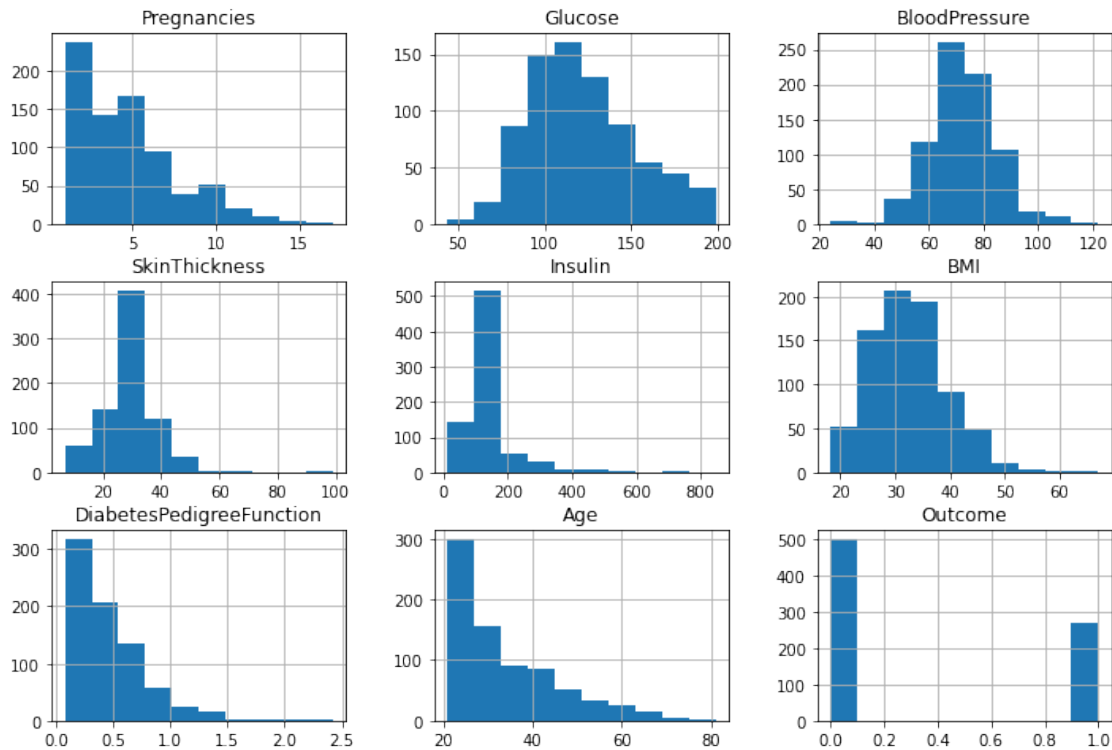
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6.000000	148.0	72.0	35.00000	155.548223	33.6	
1	1.000000	85.0	66.0	29.00000	155.548223	26.6	
2	8.000000	183.0	64.0	29.15342	155.548223	23.3	
3	1.000000	89.0	66.0	23.00000	94.000000	28.1	
4	4.494673	137.0	40.0	35.00000	168.000000	43.1	
..	
763	10.000000	101.0	76.0	48.00000	180.000000	32.9	
764	2.000000	122.0	70.0	27.00000	155.548223	36.8	
765	5.000000	121.0	72.0	23.00000	112.000000	26.2	
766	1.000000	126.0	60.0	29.15342	155.548223	30.1	
767	1.000000	93.0	70.0	31.00000	155.548223	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
[15]: df.hist(bins=10, figsize=(12, 8))
plt.suptitle('Histograms for All Columns', x=0.5, y=1.02, fontsize=16)
plt.show()
```

Histograms for All Columns



```
[55]: # We have calculated the mean for each field, by excluding the zero entries and
      ↪ replaced the zero entries with the corresponding mean.
```

```
[16]: zeros_count_per_column = df.eq(0).sum()

      print("Number of zeros in each column:")
      print(zeros_count_per_column)
```

Number of zeros in each column:

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	500
dtype:	int64

3 2. EDA

```
[17]: df.dtypes
```

```
[17]: Pregnancies          float64
      Glucose             float64
      BloodPressure       float64
      SkinThickness       float64
      Insulin             float64
      BMI                float64
      DiabetesPedigreeFunction float64
      Age                int64
      Outcome             int64
      dtype: object
```

3.1 Univariate Analysis

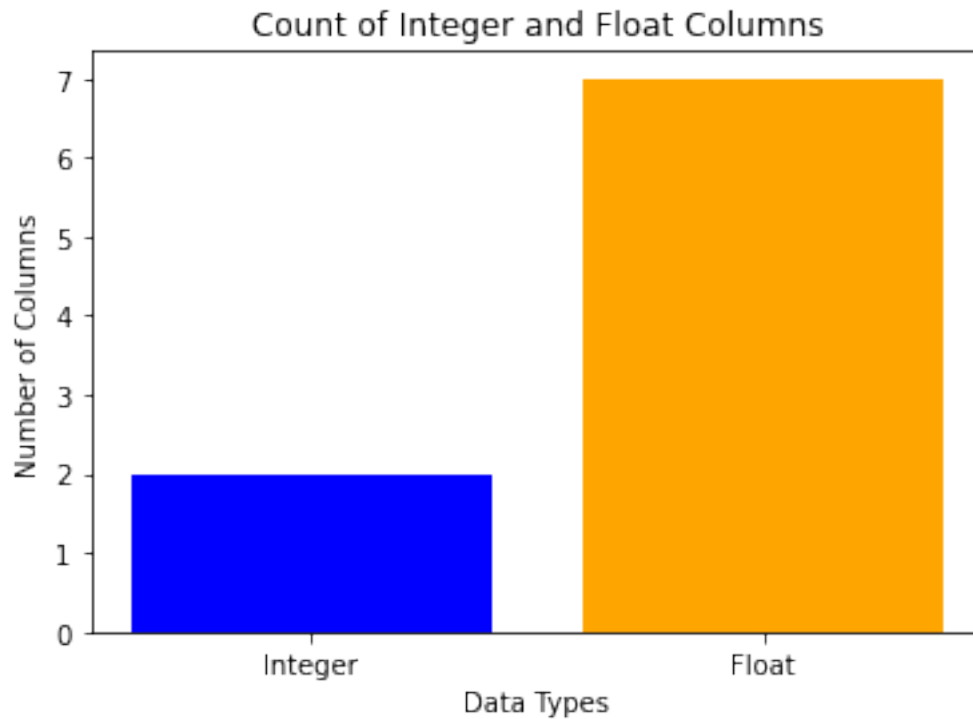
3.1.1 Creating a count (frequency) plot describing the data types and the count of variables.

```
[18]: import seaborn as sns
```

```
[19]: int_columns = df.select_dtypes(include='int').columns
      float_columns = df.select_dtypes(include='float').columns

      # Count the number of columns in each category
      int_count = len(int_columns)
      float_count = len(float_columns)

      # Plot the count plot
      plt.bar(['Integer', 'Float'], [int_count, float_count], color=['blue', 'orange'])
      plt.xlabel('Data Types')
      plt.ylabel('Number of Columns')
      plt.title('Count of Integer and Float Columns')
      plt.show()
```



[56]: *# inference : The given data set had 2 integer variables and 7 Float variables.*

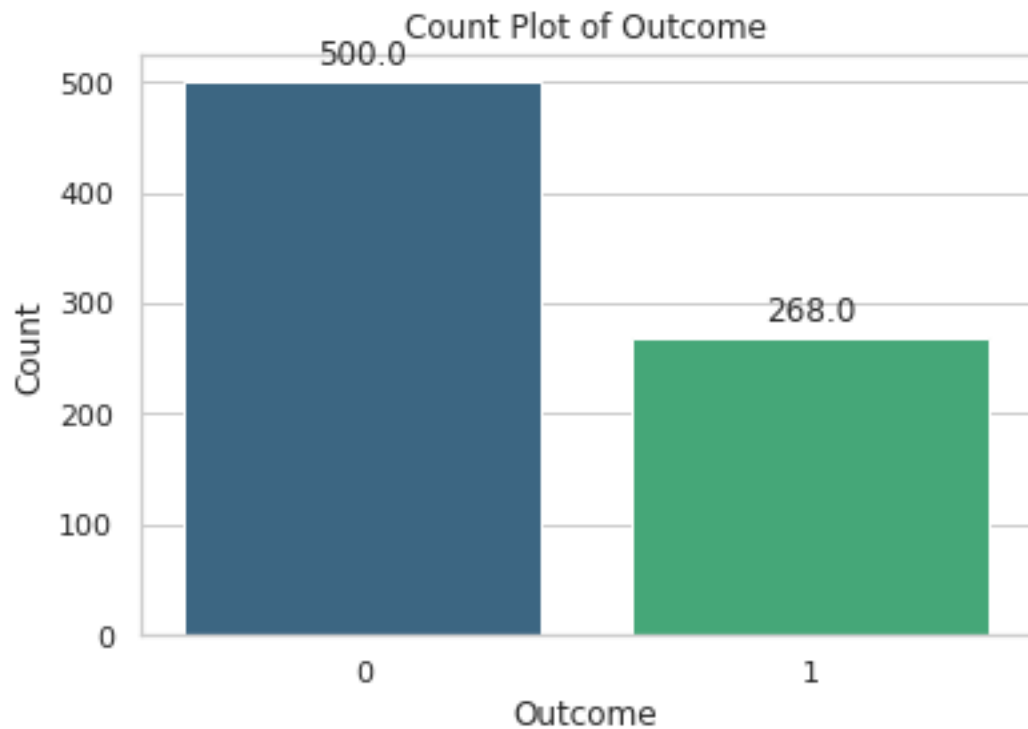
3.1.2 Plotting the count of outcomes by their value.

```
[59]: # Create a count plot for the outcome
sns.set(style="whitegrid")
plt.figure(figsize=(6, 4))
ax = sns.countplot(x="Outcome", data=df, palette='viridis')

# Add count labels on top of the bars
for p in ax.patches:
    ax.annotate(f'{p.get_height()}', (p.get_x() + p.get_width() / 2., p.
    ↪get_height()),
                ha='center', va='center', xytext=(0, 10), textcoords='offset_
    ↪points')

# Add plot labels and title
plt.xlabel('Outcome')
plt.ylabel('Count')
plt.title('Count Plot of Outcome')

# Show the plot
plt.show()
```




inference : 268 patients are positive for Diabetes against 500 negative cases. The outcome field may be a little biased towards negative outcomes.

3.2 Bivariate Analysis

3.2.1 Scatter charts between the pair of variables to understand the relationships.

```
[21]: sns.set(style="ticks")
sns.pairplot(df, hue='Outcome', markers=["o", "s"], palette='viridis')
plt.show()
```

Inference : From the histogram diagrams we can say that other than Glucose and Blood Pressure and other features have a positive skewness. Also, at a given BMI, higher glucose level is observed for diabetic patient.

3.3 Multivariate Analysis

3.3.1 Correlation analysis, using a heat map.

```
[22]: correlation_matrix = df.corr()

# Display the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)
```


Correlation Matrix:

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.154290	0.259117	0.131819	
Glucose	0.154290	1.000000	0.218367	0.192991	
BloodPressure	0.259117	0.218367	1.000000	0.192816	
SkinThickness	0.131819	0.192991	0.192816	1.000000	
Insulin	0.068077	0.420157	0.072517	0.158139	
BMI	0.110590	0.230941	0.281268	0.542398	
DiabetesPedigreeFunction	-0.005658	0.137060	-0.002763	0.100966	
Age	0.511662	0.266534	0.324595	0.127872	
Outcome	0.248263	0.492928	0.166074	0.215299	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	0.068077	0.110590	-0.005658	
Glucose	0.420157	0.230941	0.137060	
BloodPressure	0.072517	0.281268	-0.002763	
SkinThickness	0.158139	0.542398	0.100966	
Insulin	1.000000	0.166586	0.098634	
BMI	0.166586	1.000000	0.153400	
DiabetesPedigreeFunction	0.098634	0.153400	1.000000	
Age	0.136734	0.025519	0.033561	
Outcome	0.214411	0.311924	0.173844	

	Age	Outcome
Pregnancies	0.511662	0.248263
Glucose	0.266534	0.492928
BloodPressure	0.324595	0.166074
SkinThickness	0.127872	0.215299
Insulin	0.136734	0.214411
BMI	0.025519	0.311924
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[23]: plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=.5)
plt.title('Correlation Matrix Heatmap')
plt.show()
```



3.3.2 Inference : The heatmap shows that all variables have positive correlation with the target variable. blood pressue is least correlated and glucose is most correlated with target variable followed by BMI.

4 3. Data Modeling

```
[24]: import sklearn
```

```
[25]: X=df.drop('Outcome', axis = 1)
      Y=df['Outcome']
```

```
[26]: from sklearn.model_selection import train_test_split
      X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.
      ↪ 25,random_state=32)
```

```
[27]: X_train.shape
```

```
[27]: (576, 8)
```

```
[28]: X_test.shape
```

```
[28]: (192, 8)
```

5 Logistic Regression

```
[29]: from sklearn.linear_model import LogisticRegression
```

```
[30]: logmodel = LogisticRegression()
```

```
[31]: logmodel.fit(X_train,Y_train)
```

```
/usr/local/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:460:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[31]: LogisticRegression()
```

```
[32]: predictions = logmodel.predict(X_test)  
predictions
```

```
[32]: array([0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,  
        0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0,  
        0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0,  
        0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
        0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,  
        0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,  
        0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0,  
        0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0,  
        0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1])
```

```
[33]: Y_pred=logmodel.predict(X_test)  
Y_pred
```

```
[33]: array([0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
          0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0,
          0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
          0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
          0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
          0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0,
          0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0,
          0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1])
```

```
[34]: test_results=pd.DataFrame(X_test.copy())
test_results['Y_test'] = Y_test
test_results['predictions']=Y_pred
test_results
```

```
[34]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
642      6.000000    147.0      80.000000      29.15342    155.548223  29.5
695      7.000000    142.0      90.000000      24.00000    480.000000  30.4
453      2.000000    119.0      72.405184      29.15342    155.548223  19.6
565      2.000000     95.0      54.000000      14.00000     88.000000  26.1
99       1.000000    122.0      90.000000      51.00000    220.000000  49.7
..          ...      ...      ...      ...      ...      ...
585      1.000000     93.0      56.000000      11.00000    155.548223  22.5
487      4.494673    173.0      78.000000      32.00000    265.000000  46.5
58       4.494673    146.0      82.000000      29.15342    155.548223  40.5
438      1.000000     97.0      70.000000      15.00000    155.548223  18.2
394      4.000000    158.0      78.000000      29.15342    155.548223  32.9
```

	DiabetesPedigreeFunction	Age	Y_test	predictions
642	0.178	50	1	0
695	0.128	43	1	0
453	0.832	72	0	0
565	0.748	22	0	0
99	0.325	31	1	0
..
585	0.417	22	0	0
487	1.159	58	0	1
58	1.781	44	0	1
438	0.147	21	0	0
394	0.803	31	1	1

[192 rows x 10 columns]

```
[35]: from sklearn.metrics import confusion_matrix

confusion = confusion_matrix(Y_test, Y_pred)
```

```
print("Confusion Matrix:")
print(confusion)
```

Confusion Matrix:

```
[[107  16]
 [ 26  43]]
```

```
[36]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

accuracy = accuracy_score(Y_test, Y_pred)

precision = precision_score(Y_test, Y_pred)

recall = recall_score(Y_test, Y_pred)

f1 = f1_score(Y_test, Y_pred)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Accuracy: 0.78125

Precision: 0.7288135593220338

Recall: 0.6231884057971014

F1 Score: 0.671875

```
[37]: results = {'Model': [], 'Accuracy': [], 'Precision': [], 'Recall': [], 'F1 Score': []}

results['Model'].append("LogisticRegression")
results['Accuracy'].append(accuracy)
results['Precision'].append(precision)
results['Recall'].append(recall)
results['F1 Score'].append(f1)
```

5.0.1 Inference :

1. Approximately 78% of the predictions made by the model are correct.
2. A precision of 0.72 means that, out of all instances predicted as positive, 72% are truly positive.
3. A recall of 0.623 indicates that the model captures 62.3% of all actual positive instances.
4. An F1 score of 0.671 suggests a reasonable balance between precision and recall, taking into account both false positives and false negatives.

6 KNN Model

```
[38]: from sklearn.neighbors import KNeighborsClassifier
```

```
knn_clf=KNeighborsClassifier(n_neighbors=5, metric = 'euclidean')  
knn_clf.fit(X_train,Y_train)
```

```
[38]: KNeighborsClassifier(metric='euclidean')
```

```
[39]: Y_pred_knn=knn_clf.predict(X_test)  
Y_pred_knn
```

```
[39]: array([1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
        0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1,  
        0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0,  
        0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
        0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
        0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0,  
        0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0,  
        0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0,  
        0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1])
```

```
[40]: accuracy_knn = accuracy_score(Y_test, Y_pred_knn)  
  
precision_knn = precision_score(Y_test, Y_pred_knn)  
  
recall_knn = recall_score(Y_test, Y_pred_knn)  
  
f1_knn = f1_score(Y_test, Y_pred_knn)  
  
print("Accuracy:", accuracy_knn)  
print("Precision:", precision_knn)  
print("Recall:", recall_knn)  
print("F1 Score:", f1_knn)
```

```
Accuracy: 0.765625  
Precision: 0.6764705882352942  
Recall: 0.6666666666666666  
F1 Score: 0.6715328467153284
```

```
[41]: results['Model'].append("KNN")  
results['Accuracy'].append(accuracy_knn)  
results['Precision'].append(precision_knn)  
results['Recall'].append(recall_knn)  
results['F1 Score'].append(f1_knn)
```

6.0.1 Inference :

1. the model correctly predicts the class of instances in the dataset approximately 76% of the time.
2. A precision of 0.67 indicates that, out of all instances predicted as positive, around 67% are truly positive.
3. A recall of 0.66 suggests that the model captures about 66% of all actual positive instances.
4. An F1 score of 0.67 reflects a reasonable balance between precision and recall, considering both false positives and false negatives.

7 SVM Algo

```
[42]: from sklearn.svm import SVC
```

```
[43]: svm=SVC()  
      svm.fit(X_train,Y_train)
```

```
[43]: SVC()
```

```
[44]: Y_pred_svm=svm.predict(X_test)  
      Y_pred_svm
```

```
[44]: array([1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,  
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1,  
          0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,  
          0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,  
          0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,  
          0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,  
          0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0,  
          0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,  
          0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1])
```

```
[45]: accuracy_svm = accuracy_score(Y_test, Y_pred_svm)  
  
precision_svm = precision_score(Y_test, Y_pred_svm)  
  
recall_svm = recall_score(Y_test, Y_pred_svm)  
  
f1_svm = f1_score(Y_test, Y_pred_svm)  
  
print("Accuracy:", accuracy_svm)  
print("Precision:", precision_svm)  
print("Recall:", recall_svm)  
print("F1 Score:", f1_svm)
```

Accuracy: 0.8020833333333334

Precision: 0.7924528301886793

Recall: 0.6086956521739131
F1 Score: 0.6885245901639344

```
[46]: results['Model'].append("SVM")
      results['Accuracy'].append(accuracy_svm)
      results['Precision'].append(precision_svm)
      results['Recall'].append(recall_svm)
      results['F1 Score'].append(f1_svm)
```

7.0.1 Inference

1. An accuracy of approximately 80% indicates that the model correctly predicts the class of instances in the dataset the majority of the time.
2. A precision of 0.79 means that, out of all instances predicted as positive, about 79% are truly positive.
3. A recall of 0.60 indicates that the model captures about 60% of all actual positive instances.
4. An F1 score of 0.68 suggests a reasonable balance between precision and recall, considering both false positives and false negatives.

8 Decision Tree

```
[47]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score
```

```
[48]: clf = DecisionTreeClassifier()
      clf.fit(X_train, Y_train)
```

```
[48]: DecisionTreeClassifier()
```

```
[49]: Y_pred_dt = clf.predict(X_test)
      Y_pred_dt
```

```
[49]: array([0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0,
          0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0,
          0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
          0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1,
          0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1,
          0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0,
          0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0,
          0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1])
```

```
[50]: accuracy_dt = accuracy_score(Y_test, Y_pred_dt)

      precision_dt = precision_score(Y_test, Y_pred_dt)
```

```

recall_dt = recall_score(Y_test, Y_pred_dt)

f1_dt = f1_score(Y_test, Y_pred_dt)

print("Accuracy:", accuracy_dt)
print("Precision:", precision_dt)
print("Recall:", recall_dt)
print("F1 Score:", f1_dt)

```

```

Accuracy: 0.7239583333333334
Precision: 0.6025641025641025
Recall: 0.6811594202898551
F1 Score: 0.6394557823129252

```

```

[51]: results['Model'].append("Decision Tree")
      results['Accuracy'].append(accuracy_dt)
      results['Precision'].append(precision_dt)
      results['Recall'].append(recall_dt)
      results['F1 Score'].append(f1_dt)

```

8.0.1 Inference :

1. The model correctly predicts the class of instances in the dataset approximately 72% of the time.
2. A precision of 0.60 indicates that, out of all instances predicted as positive, around 60% are truly positive.
3. A recall of 0.68 suggests that the model captures about 68% of all actual positive instances.
4. An F1 score of 0.63 reflects a reasonable balance between precision and recall, considering both false positives and false negatives.

9 Comparison of Model Performances

```

[52]: # Visualize results
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 10))
fig.suptitle('Comparison of Model Performance Metrics')

metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']

for i, metric in enumerate(metrics):
    ax = axes[i // 2, i % 2]
    ax.bar(results['Model'], results[metric], color='skyblue')
    ax.set_title(metric)
    ax.set_ylim(0, 1)
    ax.set_ylabel('Score')

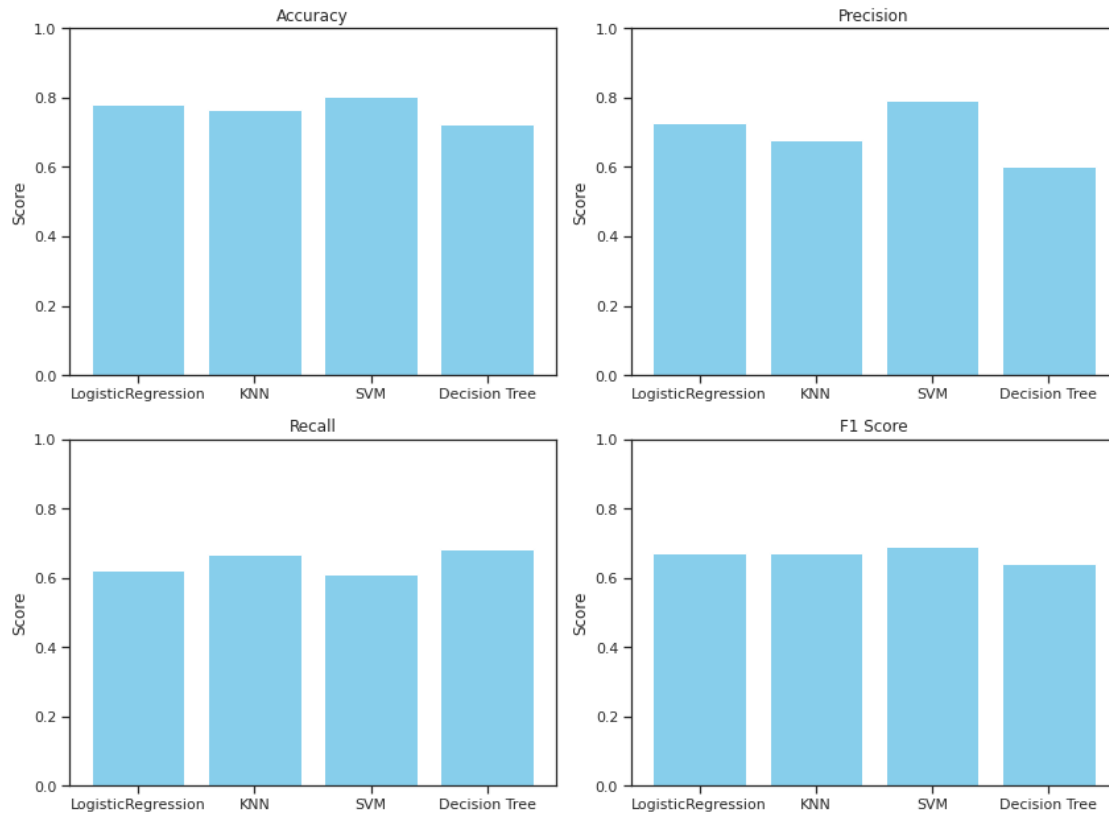
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()

```

```
# Display summary table
import pandas as pd

results_df = pd.DataFrame(results)
print(results_df)
```

Comparison of Model Performance Metrics



	Model	Accuracy	Precision	Recall	F1 Score
0	LogisticRegression	0.781250	0.728814	0.623188	0.671875
1	KNN	0.765625	0.676471	0.666667	0.671533
2	SVM	0.802083	0.792453	0.608696	0.688525
3	Decision Tree	0.723958	0.602564	0.681159	0.639456

9.0.1 Inference :

1. The SVM model has the highest accuracy, Precision and F1 Score among the four models.
2. Decision Tree has the best Recall score among the four models.

9.1 Classification report by analyzing sensitivity, specificity, AUC (ROC curve)

```
[53]: from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report, roc_auc_score, roc_curve, \
          ↪confusion_matrix

      # Create and train an SVM classifier
      svm_model = SVC(probability=True, kernel='linear')
      svm_model.fit(X_train, Y_train)

      # Make predictions and probability estimates
      Y_pred = svm_model.predict(X_test)
      Y_pred_proba = svm_model.predict_proba(X_test)[:, 1]

      # Generate classification report
      print("Classification Report:")
      print(classification_report(Y_test, Y_pred))

      # Calculate Sensitivity and Specificity
      cm = confusion_matrix(Y_test, Y_pred)
      sensitivity = cm[1, 1] / (cm[1, 0] + cm[1, 1]) # True Positive Rate or Recall
      specificity = cm[0, 0] / (cm[0, 0] + cm[0, 1])

      print(f"Sensitivity (True Positive Rate): {sensitivity:.4f}")
      print(f"Specificity (True Negative Rate): {specificity:.4f}")

      # Calculate ROC-AUC
      roc_auc = roc_auc_score(Y_test, Y_pred_proba)
      print(f"ROC-AUC: {roc_auc:.4f}")

      # Plot ROC curve
      fpr, tpr, _ = roc_curve(Y_test, Y_pred_proba)
      plt.figure(figsize=(8, 6))
      plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'AUC = {roc_auc:.2f}')
      plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('Receiver Operating Characteristic (ROC) Curve')
      plt.legend(loc='lower right')
      plt.show()
```

Classification Report:


	precision	recall	f1-score	support
0	0.79	0.89	0.84	123
1	0.74	0.58	0.65	69
accuracy			0.78	192

macro avg	0.77	0.73	0.74	192
weighted avg	0.77	0.78	0.77	192

Sensitivity (True Positive Rate): 0.5797

Specificity (True Negative Rate): 0.8862

ROC-AUC: 0.8436



9.1.1 Inference :

1. The model demonstrates good performance with high specificity (ability to correctly identify negatives) and moderate sensitivity (ability to correctly identify positives).
2. The ROC-AUC value of 0.8436 indicates strong discriminatory power, suggesting effective separation between positive and negative instances.

[]: